

	chance calibration	attribution to training	comparison model	perturbation test	matched control
F1a clinical findings and anatomy (linear)	2/18	3/18	–	–	–
F1a' the same, non-linear readouts	0/15	0/15	–	–	–
F1b held-out concept recovery, text scorers	0/6	1/6	–	–	–
F1b' concept-vector gradient association	0/2	0/2	–	–	–
F1b" labelled representational correspondence	1/5	1/5	–	–	–
F1c concept layers on foundation features	0/13	1/13	–	1/13	–
F2 sex, age or race decodable	2/9	2/9	–	–	–
F2b acquisition site and device decodable	1/12	1/12	–	–	–
F3a centre and scanner structure, pathology	0/18	0/18	3/8	–	–
F3b post-hoc harmonisation raises invariance	–	–	–	–	–
F3c site and artefact structure in brain MRI	0/3	1/3	1/2	–	–
F4 sex redundant for a refitted probe	–	–	–	–	–
F5 anisotropy and the effect of whitening	0/1	0/1	1/1	–	–
F5b medical encoders share structure	1/5	0/5	4/5	–	–
F6 modality gap; separation task-dependent	0/1	0/1	1/1	–	–
F8a attribution sets fail an instrumented test	–	1/6	–	6/6	–
F8a' attribution methods disagree	–	1/3	–	2/3	–
F8a" maps implausible on inspection	–	0/2	–	0/2	–
F8b labels less decodable after the connector	–	0/1	–	0/1	–
F9a targeted interventions report changed outputs	–	–	–	–	3/8
F7 reported validation of dictionary features	row selected on semantic validation; no share shown				
F9b inventory: editing and unlearning	–	–	–	–	0/1
F10 inventory: medical vs general decodability	1/20	2/20	–	0/3	–
F11 inventory: properties shared or differing	4/19	3/21	11/11	2/2	–

controls reported by the row's records, as k/n over the records whose design can carry the control